



Application of Deep Learning to Text and Image Data

Module 1, Lab 3: Building an End-to-End Neural Network Solution

In the previous lab, you used a neural network with image data to predict the category that an item belonged to. In this lab, you will process text data by building an end-to-end neural network solution. The solution will incorporate all the data processing techniques that you have learned so far.

You will learn how to do the following:

- Import and preprocess data.
- Create a neural network with multiple layers.
- Train text data with your neural network.
- Validate your model as you train.
- Change different parameters to improve your neural network.

Austin Animal Center Dataset

In this lab, you will work with historical pet adoption data in the [Austin Animal Center Shelter Intakes and Outcomes dataset](https://www.kaggle.com/datasets/aaronschlegel/austin-animal-center-shelter-intakes-and-outcomes?resource=download) (<https://www.kaggle.com/datasets/aaronschlegel/austin-animal-center-shelter-intakes-and-outcomes?resource=download>). The target field of the dataset (**Outcome Type**) is the outcome of adoption: 1 for adopted and 0 for not adopted. Multiple features are used in the dataset.

Dataset schema:

- **Pet ID:** Unique ID of the pet
- **Outcome Type:** State of pet at the time of recording the outcome (0 = not placed, 1 = placed). This is the field to predict.

- **Sex upon Outcome:** Sex of pet at outcome
 - **Name:** Name of pet
 - **Found Location:** Found location of pet before it entered the shelter
 - **Intake Type:** Circumstances that brought the pet to the shelter
 - **Intake Condition:** Health condition of the pet when it entered the shelter
 - **Pet Type:** Type of pet
 - **Sex upon Intake:** Sex of pet when it entered the shelter
 - **Breed:** Breed of pet
 - **Color:** Color of pet
 - **Age upon Intake Days:** Age (days) of pet when it entered the shelter
 - **Age upon Outcome Days:** Age (days) of pet at outcome
-

You will be presented with two kinds of exercises throughout the notebook: activities and challenges.



No coding is needed for an activity. You try to understand a concept, answer questions, or run a code cell.



Challenges are where you can practice your coding skills.

Index

- [Data processing](#)
 - [Training and validation of a neural network](#)
 - [Testing the neural network](#)
 - [Improvement ideas](#)
-

Data processing

The first step is to process the dataset.

```
In [17]: # Install Libraries
!pip install -U -q -r requirements.txt
```

```
In [18]: # Import the dependencies
import boto3
import os
from os import path
import pandas as pd

%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import re, string
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, MinMaxScaler
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.utils import shuffle
import torch
from torch import nn
import Stemmer

from MLUDTI_M1_Lab3_neural_network import NeuralNetwork
```

First, read the dataset into a DataFrame and look at it. The data might look familiar because it was used in the labs of the Machine Learning through Application course.

```
In [19]: df = pd.read_csv("data/austin-animal-center-dataset.csv")

print("The shape of the dataset is:", df.shape)
```

The shape of the dataset is: (95485, 13)

```
In [20]: # Print the first five rows of the dataset
df.head()
```

Out[20]:

	Pet ID	Outcome Type	Sex upon Outcome	Name	Found Location	Intake Type	Intake Condition	Pet Type	Sex upon Intake
0	A794011	1.0	Neutered Male	Chunk	Austin (TX)	Owner Surrender	Normal	Cat	Neutered Male
1	A776359	1.0	Neutered Male	Gizmo	7201 Levander Loop in Austin (TX)	Stray	Normal	Dog	Intact Male
2	A674754	0.0	Intact Male	NaN	12034 Research in Austin (TX)	Stray	Nursing	Cat	Intact Male
3	A689724	1.0	Neutered Male	*Donatello	2300 Waterway Bnd in Austin (TX)	Stray	Normal	Cat	Intact Male
4	A680969	1.0	Neutered Male	*Zeus	4701 Staggerbrush Rd in Austin (TX)	Stray	Nursing	Cat	Intact Male

EDA

Now, perform the basic steps of exploratory data analysis (EDA) and look for insights to inform later ML modeling choices.

```
In [5]: # Print the data types and nonnull values for each column
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 95485 entries, 0 to 95484
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pet ID                95485 non-null  object
1   Outcome Type          95485 non-null  float64
2   Sex upon Outcome      95484 non-null  object
3   Name                  59138 non-null  object
4   Found Location        95485 non-null  object
5   Intake Type           95485 non-null  object
6   Intake Condition      95485 non-null  object
7   Pet Type              95485 non-null  object
8   Sex upon Intake       95484 non-null  object
9   Breed                 95485 non-null  object
10  Color                 95485 non-null  object
11  Age upon Intake Days  95485 non-null  int64
12  Age upon Outcome Days 95485 non-null  int64
dtypes: float64(1), int64(2), object(10)
memory usage: 9.5+ MB
```

```
In [6]: # Print the column names
print(df.columns)
```

```
Index(['Pet ID', 'Outcome Type', 'Sex upon Outcome', 'Name', 'Found Location',
      'Intake Type', 'Intake Condition', 'Pet Type', 'Sex upon Intake',
      'Breed', 'Color', 'Age upon Intake Days', 'Age upon Outcome Days'],
      dtype='object')
```

```
In [7]: # Create lists that identify the numerical, categorical, and text features, and
numerical_features = ["Age upon Intake Days", "Age upon Outcome Days"]
```

```
categorical_features = [
    "Sex upon Outcome",
    "Intake Type",
    "Intake Condition",
    "Pet Type",
    "Sex upon Intake",
]
```

```
text_features = ["Found Location", "Breed", "Color"]
```

```
model_features = numerical_features + categorical_features + text_features
model_target = "Outcome Type"
```

Note: The **Pet ID** and **Name** features were omitted because they are irrelevant to the outcome.

Cleaning the data

Cleaning numerical features

Take a moment to examine the numerical features. Remember that the `value_counts()` function can give a view of the numerical features by placing feature values in respective bins. The function can also be used for plotting.


```
In [8]: for c in numerical_features:
        print(c)
        print(df[c].value_counts(bins=10, sort=False))
        df[c].value_counts(bins=10, sort=False).plot(kind="bar", alpha=0.75, rot=90)
        plt.show()
```

Age upon Intake Days


Age upon Intake Days

(-9.126, 912.5]	74835
(912.5, 1825.0]	10647
(1825.0, 2737.5]	3471
(2737.5, 3650.0]	3998
(3650.0, 4562.5]	1234
(4562.5, 5475.0]	1031
(5475.0, 6387.5]	183
(6387.5, 7300.0]	79
(7300.0, 8212.5]	5
(8212.5, 9125.0]	2

Name: count, dtype: int64



```
Age upon Outcome Days
Age upon Outcome Days
(-9.126, 912.5]      74642
(912.5, 1825.0]     10699
(1825.0, 2737.5]     3465
(2737.5, 3650.0]     4080
(3650.0, 4562.5]     1263
(4562.5, 5475.0]     1061
(5475.0, 6387.5]      187
(6387.5, 7300.0]       81
(7300.0, 8212.5]        5
(8212.5, 9125.0]        2
Name: count, dtype: int64
```



If any outliers are identified as likely wrong values, dropping them could improve the histograms for the numerical values and could later improve overall model performance.

Remove any values in the upper 10 percent for the feature, and then plot the features.


```
In [9]: for c in numerical_features:
        # Drop values beyond 90% of max()
        dropIndexes = df[df[c] > df[c].max() * 9 / 10].index
        df.drop(dropIndexes, inplace=True)
```

```
In [10]: for c in numerical_features:
          print(c)
          print(df[c].value_counts(bins=10, sort=False))
          df[c].value_counts(bins=10, sort=False).plot(kind="bar", alpha=0.75, rot=90)
          plt.show()
```

Age upon Intake Days

Age upon Intake Days

(-6.936, 693.5] 61425

(693.5, 1387.0] 18400

(1387.0, 2080.5] 5657

(2080.5, 2774.0] 3471

(2774.0, 3467.5] 2557

(3467.5, 4161.0] 1962


(4161.0, 4854.5] 1148

(4854.5, 5548.0] 596


(5548.0, 6241.5] 183

(6241.5, 6935.0] 63

Name: count, dtype: int64



```
Age upon Outcome Days
Age upon Outcome Days
(-6.936, 693.5]      61208
(693.5, 1387.0]     18490
(1387.0, 2080.5]     5643
(2080.5, 2774.0]     3465
(2774.0, 3467.5]     2600
(3467.5, 4161.0]     2004
(4161.0, 4854.5]     1196
(4854.5, 5548.0]      604
(5548.0, 6241.5]      187
(6241.5, 6935.0]       65
Name: count, dtype: int64
```



Cleaning text features

Take a moment to examine the text features.

```
In [11]: # Prepare cleaning functions
import re, string

stop_words = ["a", "an", "the", "this", "that", "is", "it", "to", "and", "in"]

stemmer = Stemmer.Stemmer('english')

def preProcessText(text):
    # Lowercase text, and strip leading and trailing white space
    text = text.lower().strip()

    # Remove HTML tags
    text = re.compile("<.*?>").sub("", text)

    # Remove punctuation
    text = re.compile("[%s]" % re.escape(string.punctuation)).sub(" ", text)

    # Remove extra white space
    text = re.sub("\s+", " ", text)

    return text

def lexiconProcess(text, stop_words, stemmer):
    filtered_sentence = []
    words = text.split(" ")
    for w in words:
        if w not in stop_words:
            filtered_sentence.append(stemmer.stemWord(w))
    text = " ".join(filtered_sentence)

    return text

def cleanSentence(text, stop_words, stemmer):
    return lexiconProcess(preProcessText(text), stop_words, stemmer)
```

Note: The text cleaning process can take a while to complete, depending on the size of the text data.

```
In [12]: # Clean the text features
for c in text_features:
    print("Text cleaning: ", c)
    df[c] = [cleanSentence(item, stop_words, stemmer) for item in df[c].values]
```

```
Text cleaning: Found Location
Text cleaning: Breed
Text cleaning: Color
```

Train, validation, and test datasets

Now that the data has been cleaned, you need to split the full dataset into training and test subsets by using sklearn's `train_test_split()` (https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html) function. With

this function, you can specify the following:

- The proportion of the dataset to include in the test split as a number between 0.0-1.0 with a default of 0.25.
- An integer that controls the shuffling that is applied to the data before the split. Passing an integer allows for reproducible output across multiple function calls.

To help reduce sampling bias, the original dataset is shuffled before the split. After the initial split, the training data is further split into training and validation subsets.

```
In [13]: from sklearn.model_selection import train_test_split

train_data, test_data = train_test_split(
    df, test_size=0.15, shuffle=True, random_state=23
)

train_data, val_data = train_test_split(
    train_data, test_size=0.15, shuffle=True, random_state=23
)

# Print the shapes of the training, validation, and test datasets
print(
    "Train - Validation - Test dataset shapes: ",
    train_data.shape,
    val_data.shape,
    test_data.shape,
)
```

```
Train - Validation - Test dataset shapes: (68970, 13) (12172, 13) (14320, 13)
```

Data processing with a pipeline and ColumnTransformer

In a typical ML workflow, you need to apply data transformations, such as imputation and scaling, at least twice: first on the training dataset by using `.fit()` and `.transform()` when preparing the data to train the model, and then by using `.transform()` on any new data that you want to predict on (validation or test). Sklearn's [Pipeline \(https://scikit-learn.org/stable/modules/generated/sklearn.pipeline.Pipeline.html\)](https://scikit-learn.org/stable/modules/generated/sklearn.pipeline.Pipeline.html) is a tool that simplifies this process by enforcing the implementation and order of data processing steps.

In this section, you will build separate pipelines to handle the numerical, categorical, and text features. Then, you will combine them into a composite pipeline along with an estimator. To do this, you will use a [LogisticRegression classifier \(https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html\)](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html).

You will need multiple pipelines to ensure that all the data is handled correctly:

- **Numerical features pipeline:** Impute missing values with the mean by using sklearn's `SimpleImputer`, followed by `MinMaxScaler`. If different processing is desired for different numerical features, different pipelines should be built as described for the text features pipeline.

- **Categoricals pipeline:** Impute with a placeholder value (this won't have an effect because you already encoded the `nan` values), and encode with sklearn's `OneHotEncoder`. If computing memory is an issue, it is a good idea to check the number of unique values for the categoricals to get an estimate of how many dummy features one-hot encoding will create. Note the `handle_unknown` parameter, which tells the encoder to ignore (rather than throw an error for) any unique value that might show in the validation or test set that was not present in the initial training set.
- **Text features pipeline:** With memory usage in mind, build three more pipelines, one for each of the text features. The current sklearn implementation requires a separate transformer for each text feature (unlike the numericals and categoricals).

Finally, the selective preparations of the dataset features are then put together into a collective `ColumnTransformer`, which is used in a pipeline along with an estimator. This ensures that the transforms are performed automatically in all situations. This includes on the raw data when fitting the model, when making predictions, when evaluating the model on a validation dataset


```

In [14]: from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, MinMaxScaler
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer

### COLUMN_TRANSFORMER ###
#####

# Preprocess the numerical features
numerical_processor = Pipeline(
    [
        ("num_imputer", SimpleImputer(strategy="mean")),
        (
            "num_scaler",
            MinMaxScaler(),
        ),
    ]
)

# Preprocess the categorical features
categorical_processor = Pipeline(
    [
        (
            "cat_imputer",
            SimpleImputer(strategy="constant", fill_value="missing"),
        ), # Shown in case it is needed. No effect here because you already impute
        (
            "cat_encoder",
            OneHotEncoder(handle_unknown="ignore"),
        ), # handle_unknown tells it to ignore (rather than throw an error for
    ]
)

# Preprocess first text feature
text_processor_0 = Pipeline(
    [("text_vectorizer_0", CountVectorizer(binary=True, max_features=50))]
)

# Preprocess second text feature
text_processor_1 = Pipeline(
    [("text_vectorizer_1", CountVectorizer(binary=True, max_features=50))]
)

# Preprocess third text feature
text_processor_2 = Pipeline(
    [("text_vectorizer_2", CountVectorizer(binary=True, max_features=50))]
)

# Combine all data preprocessors (add more if you choose to define more)
# For each processor/step, specify: a name, the actual process, and the features
data_processor = ColumnTransformer(
    [
        ("numerical_processing", numerical_processor, numerical_features),
        ("categorical_processing", categorical_processor, categorical_features),
        ("text_processing_0", text_processor_0, text_features[0]),
        ("text_processing_1", text_processor_1, text_features[1]),
    ]
)

```



```

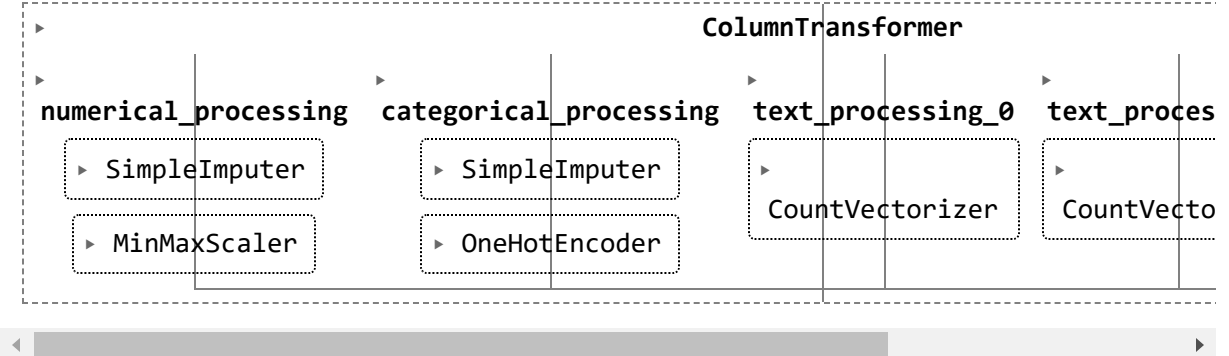
        ("text_processing_2", text_processor_2, text_features[2]),
    ]
)

# Visualize the data processing pipeline
from sklearn import set_config

set_config(display="diagram")
data_processor

```

Out[14]:



```

In [15]: # Prepare data for training
X_train = train_data[model_features]
y_train = train_data[model_target].values

# Get validation data to validate the network
X_val = val_data[model_features]
y_val = val_data[model_target].values

# Get test data to test the network for submission to the leaderboard
X_test = test_data[model_features]
y_test = test_data[model_target].values

print("Dataset shapes before processing: ", X_train.shape, X_val.shape, X_test.shape)

X_train = data_processor.fit_transform(X_train).toarray()
X_val = data_processor.transform(X_val).toarray()
X_test = data_processor.transform(X_test).toarray()

print("Dataset shapes after processing: ", X_train.shape, X_val.shape, X_test.shape)

```

```

Dataset shapes before processing: (68970, 10) (12172, 10) (14320, 10)
Dataset shapes after processing: (68970, 171) (12172, 171) (14320, 171)

```

Training and validation of a neural network

Now, run the following code cell to interact with the neural network to gain insight into how neural networks train.

Architect the neural network by updating the number of layers (maximum of 4) and the number of neurons per layer (maximum of 3) to solve the classification problem that displays when you run the following cell. The background colors show the neural network's predicted classification

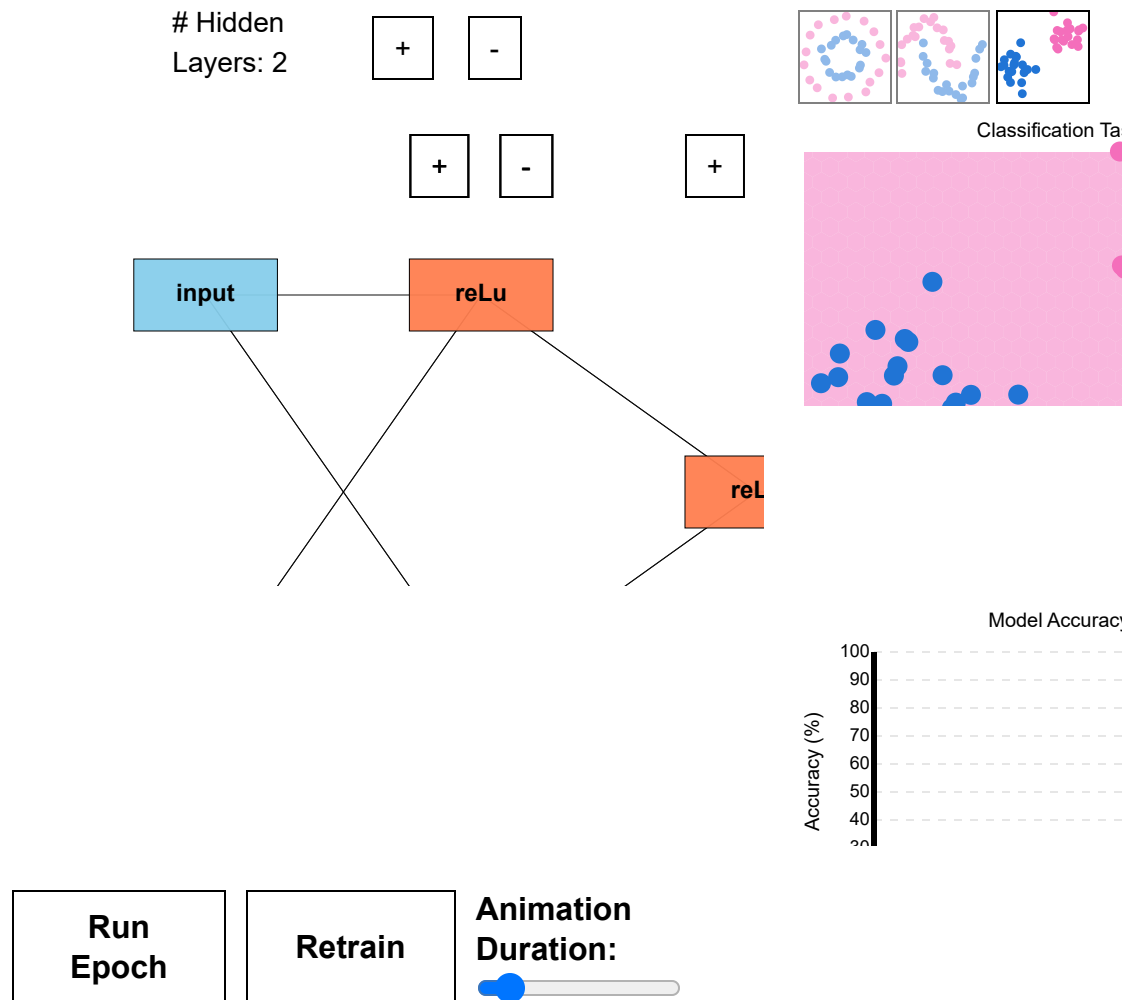
regions for the true data (circles).

Note that upon retraining the network, the weights are randomly initialized, and the gradients are reset to 0. In the visual representation, each green circle corresponds to an epoch. Each red circle corresponds to that layer's weight update gradient (from backpropagation).

To develop a better understanding, train the model for different architectures. Note that the model gets stuck sometimes—initialization is important!

In [16]: `NeuralNetwork()`

Out[16]:



Now you need to build a PyTorch neural network and use it to fit to the training data. As part of the training, you need to use the validation data to check performance at the end of each training iteration.

Try it yourself!



To define the hyperparameters of the algorithm, run the following cell. Note how the data is loaded into PyTorch tensors. Observe how the DataLoader is defined. The DataLoader is used

```
In [21]: # Define the hyperparameters
batch_size = 16
num_epochs = 15
learning_rate = 0.001
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

# Convert the data into PyTorch tensors

X_train = torch.tensor(X_train, dtype=torch.float32).to(device)
X_val = torch.tensor(X_val, dtype=torch.float32).to(device)
X_test = torch.tensor(X_test, dtype=torch.float32).to(device)

y_train = torch.tensor(y_train, dtype=torch.long).to(device)
y_val = torch.tensor(y_val, dtype=torch.long).to(device)
y_test = torch.tensor(y_test, dtype=torch.long).to(device)

# Use PyTorch DataLoaders to load the data in batches
train_dataset = torch.utils.data.TensorDataset(X_train, y_train)
train_loader = torch.utils.data.DataLoader(train_dataset,
                                           batch_size=batch_size,
                                           drop_last=True)

val_dataset = torch.utils.data.TensorDataset(X_val, y_val)
val_loader = torch.utils.data.DataLoader(val_dataset,
                                         batch_size=batch_size,
                                         drop_last=True)
```

Try it yourself!



Challenge

Create a multilayer perceptron by using the `Sequential` module with the following attributes:

1. Use two hidden layers, both of size 64.
2. Attach a dropout layer to each hidden layer.
3. Use a ReLU activation for each hidden layer.
4. Create one output layer.

```
In [31]: # Create a multilayer perceptron by using the Sequential module. Add the follow
# Two hidden layers of size 64
# Dropout layers attached to the hidden layers
# ReLU activation functions
# One output layer

##### CODE HERE #####
net = nn.Sequential(
    nn.Linear(in_features=171, out_features=64),
    nn.ReLU(),
    nn.Dropout(p=0.3),

    nn.Linear(in_features=64, out_features=64),
    nn.ReLU(),
    nn.Dropout(p=0.3),

    nn.Linear(in_features=64, out_features=10) # Assume 10 classes for classification
).to(device)
##### END OF CODE #####

def xavier_init_weights(m):
    if type(m) == nn.Linear:
        torch.nn.init.xavier_uniform_(m.weight)

net.apply(xavier_init_weights)
```

```
Out[31]: Sequential(
  (0): Linear(in_features=171, out_features=64, bias=True)
  (1): ReLU()
  (2): Dropout(p=0.3, inplace=False)
  (3): Linear(in_features=64, out_features=64, bias=True)
  (4): ReLU()
  (5): Dropout(p=0.3, inplace=False)
  (6): Linear(in_features=64, out_features=10, bias=True)
)
```

```
In [32]: # Define the loss function and the optimizer
# Choose cross-entropy loss for this classification problem
loss = nn.CrossEntropyLoss()

# Optimize with stochastic gradient descent. You can experiment with other optimizers
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate)
```

Try it yourself!



Activity

To train the network, run the following cell. Watch how the training and validation loss change for each epoch.

```

In [33]: import time

#####
# Network training and validation

# Start the outer epoch loop (epoch = full pass through the dataset)
for epoch in range(num_epochs):
    start = time.time()

    training_loss, validation_loss = 0.0, 0.0

    # Training Loop (with autograd and trainer steps)
    # This Loop trains the neural network
    # Weights are updated here
    net.train() # Activate training mode (dropouts and so on)
    for data, target in train_loader:
        # Zero the parameter gradients
        optimizer.zero_grad()
        data = data.to(device)
        target = target.to(device)
        # Forward + backward + optimize
        output = net(data)
        L = loss(output, target)
        L.backward()
        optimizer.step()
        # Add batch Loss
        training_loss += L.item()

    net.eval() # Activate eval mode (don't use dropouts and so on)
    for data, target in val_loader:
        data = data.to(device)
        target = target.to(device)
        output = net(data)
        L = loss(output, target)
        # Add batch Loss
        validation_loss += L.item()

    # Take the average Losses
    training_loss = training_loss / len(train_loader)
    val_loss = validation_loss / len(val_loader)

    end = time.time()
    print(
        "Epoch %s. Train_loss %f Validation_loss %f Seconds %f"
        % (epoch, training_loss, val_loss, end - start)
    )

```

```
Epoch 0. Train_loss 0.934593 Validation_loss 0.605567 Seconds 14.224415
Epoch 1. Train_loss 0.613182 Validation_loss 0.521049 Seconds 4.405684
Epoch 2. Train_loss 0.548573 Validation_loss 0.470559 Seconds 4.396922
Epoch 3. Train_loss 0.507711 Validation_loss 0.450232 Seconds 4.388942
Epoch 4. Train_loss 0.486486 Validation_loss 0.440889 Seconds 4.345073
Epoch 5. Train_loss 0.476056 Validation_loss 0.435685 Seconds 4.418920
Epoch 6. Train_loss 0.466256 Validation_loss 0.432310 Seconds 4.414241
Epoch 7. Train_loss 0.459010 Validation_loss 0.429688 Seconds 4.451204
Epoch 8. Train_loss 0.455062 Validation_loss 0.427885 Seconds 4.574066
Epoch 9. Train_loss 0.450703 Validation_loss 0.426692 Seconds 4.417195
Epoch 10. Train_loss 0.448325 Validation_loss 0.425762 Seconds 4.395355
Epoch 11. Train_loss 0.446277 Validation_loss 0.424805 Seconds 4.410755
Epoch 12. Train_loss 0.443940 Validation_loss 0.423619 Seconds 4.398337
Epoch 13. Train_loss 0.440316 Validation_loss 0.423132 Seconds 4.390956
Epoch 14. Train_loss 0.438620 Validation_loss 0.422772 Seconds 4.528893
```

Testing the neural network

Now you can evaluate the performance of the trained network on the test set.

```
In [34]: from sklearn.metrics import classification_report

# Activate eval mode (don't use dropouts and so on)
net.eval()

# Get test predictions
predictions = net(X_test)

# Print performance of the test data
print(
    classification_report(
        y_test.cpu().numpy(), predictions.argmax(axis=1).cpu().detach().numpy()
    )
)
```

	precision	recall	f1-score	support
0	0.86	0.71	0.78	6267
1	0.80	0.91	0.85	8053
accuracy			0.83	14320
macro avg	0.83	0.81	0.82	14320
weighted avg	0.83	0.83	0.82	14320

Improvement ideas

You can improve this neural network by tuning network parameters such as the following:

- Architecture
- Number of layers
- Number of hidden neurons
- Choice of activation function
- Weight initialization
- Dropout
- Choice of optimizer function
- Learning rate
- Batch size
- Number of epochs

As you make changes, closely monitor the loss function and the accuracy on both training and validation to identify what changes improve your model.

Conclusion

In this notebook, you built a basic neural network to process text data.

Next Lab: Introducing CNNs

In the next lab in this module you will learn how to build a convolutional neural network to process hand written numbers.

In []: